

Montalbán working-paper style: everything learned

Compiled 28-09-2026 from style/our_skill/montalban-wp/SKILL.md and style/blindtest/LOG.md

Montalbán working-paper style, v2 (consolidated 24-09-2026)

A. Corpus and evidence

The corpus is in C:_repl_jornada_paper\style\montalban\ (INDEX.md). The measured units are in units/, and the matrix is in _evidence_matrix.txt, produced by tools/corpus_features.py.

Evidence units, one per paper, each in its latest WP version:

Unit	Paper	Year	Authors	Madrid data
AA2025	Affirmative Action during Early Childhood	2025	Margaryan & Montalbán	LOMCE + questionnaires
BB2024	Home Broadband (CEP DP)	2024	Montalbán, Sanchis-Guarner, Weinhardt	no
SC2020may	School Choice Priorities (SOFI WP)	2020	Gortázar, Mayor, Montalbán	CDI as a control
JMP2019nov	Countering Moral Hazard (PSE WP v2)	2019	sole	no

GENDER2019 (thesis ch. 2) is EXCLUDED from the evidence (user, 24-09-2026: a PhD thesis is a different genre from a WP). It is kept only as background on Madrid testing data. The PhD thesis as a whole is also excluded. Evolution units, used only to see how the style changed: BB2021 (QMUL), SC2020jan, JMP2019jun, JMP2018.

Weighting rules (user):

- All his WPs count as his voice, since he was usually the junior author and did the writing.
- A paper gets less weight only when the co-authors are very senior. Pipeline vs Choice is excluded. "Smaller Classes" (Clark, Deschamps) would count fully, but it is not available.
- Underdeveloped early drafts get less weight.

Evidence levels used below:

- [4/4], [3/4], ... = number of evidence units in which the rule holds.
- **STABLE** = at least 3/4 and present in at least one of the 2024–25 papers → apply.
- **ERA-2019/20** = present in SC and JMP but absent from AA2025 and BB2024 → apply sparingly (default Q1, section G).
- **SINGLE** = one paper only → do not use as a rule.

Anti-AI layer: style/our_skill/anti_ai_patterns.md (Lu Han, verbatim). **When it conflicts with Montalbán, Montalbán wins** (user decision). No other style guide applies.

B. Process (from the user; comes before style)

1. Ideas before prose.

- Before writing a section, draft an argument plan: one line per paragraph saying what it *defends* and how it leads to the next.
 - List the **gaps**: why something matters, what to compare it with, what the author would claim.
2. **Author's questions go to the user before writing.** These include benchmarks, the preferred estimate, how strongly to interpret, what goes in the text versus a footnote or appendix, and any institutional fact not in our sources. Do not write, find the problem, then patch.
 3. **Numbers** are phrased as in section E. Checking them against facts/ is a silent background step.
 4. **Never invent facts or interpretations** that the tables and facts do not give. Round 1 showed the temptation, e.g. "roughly what we would expect by chance".
 5. **Empirical Strategy** is presented intuitively and plainly: the comparison in words, why it identifies the effect, then the equation. It should not read like an econometrician showing off.

6. Getting close is enough; 100% is not expected.

C. Architecture and lengths

Measure	AA2025	BB2024	SC2020	JMP2019	(GENDER, thesis, excluded)	Rule
Main-text words	7,488	12,406	11,029	13,020	7,129	7.5–13k. Target ~9–1 for Jornada
Paragraph median	104	128	150	147	121	STABLE 100–150
Paragraph p10 / p90	52/177	57/231	62/281	54/310	68/270	STABLE short ones ~55, long ones 180–3
Sentence median (words)	20	22	25	26	23	STABLE 20–26, with large spread (SD 10–)
Footnotes per 1k words	1.7	2.7	4.3	3.0	1.1	STABLE 1–4, mean 33–61 words each
Intro words / paragraphs	1,379/11	1,638/10	1,847/10	2,714/12	1,681/11	STABLE [4/4] 10–12 paragraphs, 1,400–2, words

- Numbering: "1 Introduction", "2.1 ..." (arabic). [4/4]
- **Run-in headings** inside sections. Evidence: SC and JMP use Title Case ("Sample Restrictions.", "Average Effects.", "Eligibility Rules."); AA uses sentence case ("Household school choice patterns."); BB uses numbered subsections instead. [3/4: AA, SC, JMP] → use them, in sentence case (default Q2, section G).
- Paragraph lengths vary. There are no one-sentence punch paragraphs: p10 is never below ~50 words. [4/4]

D. Section templates

Introduction (arc confirmed across units; paragraph numbers from the matrix)

- P1: the policy problem and why it matters, with 2–4 citations. It may pack motivation, the trade-off and a budget or scale figure into one long paragraph. [4/4] have the densest citation block at P1 or at the end.
- "This paper / In this paper / We study ..." arrives at **P2** (AA, JMP, SC-jan) or **P4–5** (SC-may, BB). **STABLE: by P2–P5; never after the literature review.**
- The headline results with numbers come early. They appear in P1–P2 in JMP and AA, and in P5–P6 elsewhere, and come with their relative size.
- The contributions and the broad literature come **late**: the paragraph with the most citations is P8–P10 in AA, SC, JMP and BB2021. [3/4] (BB2024 is the exception). The *form* varies:
 - AA: "Our study contributes to two strands of literature"
 - JMP: the run-in "Related literature."
 - SC: a literature paragraph

Rule: survey the literature after the results, not as an opening "gap".

- Roadmap at the end ("The remainder of the paper is organized as follows" / "Organization of the Paper."). [4/4]
- **Avoid** (round 2): symmetric double scaffolds ("There are two main reasons ... First ... Second ..." followed by "This setting offers two advantages. First ... Second ..."). First/Second is fine inside a paragraph; the whole intro must not be built on parallel pairs.

Institutional Background

- Go from general to specific, with concrete magnitudes and their sources. The law or decree goes in a footnote, and so do comparisons that give scale. [4/4]
- Procedures are walked through in time order (calendar). Bullets are acceptable for a dated calendar or a rule schedule (JMP, twice). Otherwise use prose.
- Close by linking the institution to identification (what it lets us compare). [4/4]

Data

- One paragraph per source (who produces it, years, universe, contents, stakes), or a list "(i) ... (ii) ..." followed by one sentence per file. [4/4] for the per-source description.
- Sample restrictions: each one with its **share** and its **reason**. Consequences for identification go in a footnote ("available upon request" or an appendix table). [4/4]
- Summary statistics walked through in prose, mentioning 3–5 facts, not the whole table. [4/4]
- Survey indices described in words: the items, the scale, the standardization. [AA]; AA is the only WP with questionnaires.

Empirical Strategy

- The comparison in words first, then "Our equation of interest ..." or "The reduced-form equation ...", then a "where ..." clause that defines each term. [4/4]
- Estimation choices (kernel, bandwidth, clustering, software) take one short sentence each, and each alternative goes in a footnote. [3/4]
- "Let X denote" before an equation: **SINGLE (JMP)** → optional.

Internal Validity

- Identifying assumption → testable implications → one paragraph per test (figure or table, result in words, complementary test in the appendix). [AA, BB, JMP]
- Balance and null tests are **terse in the text**: no covariate-by-covariate detail and no chi2 or p-values in the prose, just a pointer to the table. [4/4]; the round-1 judge flagged my extra numbers.

Results

- Descriptive picture or baseline levels first, then the estimates, then the interpretation. [4/4]
- The literature enters Results to interpret the magnitudes ("consistent with", "in line with", "contrasts with"). [4/4]
- Robustness is one sentence in the text plus a footnote or appendix table. A robustness *section* has one titled run-in paragraph per check, in the order announced. [SC, JMP, BB]
- Heterogeneity: list the values for the groups, then compare absolute versus relative size. [SC, JMP, AA]

Conclusions

- 1–6 paragraphs: restate the findings in words, then the implication, then the limits and open questions. [4/4]

E. How numbers and results are stated

- **STABLE** [4/4]:
 - Effect in natural units, then its size relative to the mean: "..., which corresponds to ...", "(X percent with respect to the baseline mean)", "or X percent relative to a control mean of Y". The exact phrasing varies by paper.
 - Significance in words ("statistically significant at the 5% level", "close to zero and statistically insignificant"). Standard errors and CIs stay in the tables.
 - "precisely estimated null", "we cannot reject ...", "we fail to reject ..." for nulls.
 - Point to the exact location: "column 5 of Panel B in Table 3", "Panels (c) and (d) of Figure 5".
 - Scale words: about, approximately, roughly, around.
- **"respectively"** to pair values: [3/4] (not AA).
- Walking through the columns and naming the preferred specification: [BB, SC, JMP].
- **Paired values in parentheses** ("65 (90) percent for STEM (non-STEM)"): **ERA-2019/20** [2/4: SC, JMP]. Absent from AA and BB. My round-1 [S] tag was **WRONG**.
- Percent: the corpus mixes "percent" and "%". **Open question Q3**.

F. Voice and connectors (rates per 1,000 words in `_evidence_matrix.txt`)

- **STABLE:**
 - “we”, active voice.
 - First/Second/Finally within paragraphs (0.7–2.4 per 1k). [4/4]
 - “However” (0.4–1.5). [4/4]
 - “In addition”, “Moreover”. [4/4]
 - “e.g.”, “i.e.”. [4/4]
 - “robust”. [4/4]
 - “potential(ly)” (0.9–1.8). [4/4]
 - “may/might” (1.1–4.7). [4/4]
 - “consistent with”. [4/4]
 - “Overall,”. [3/4]
 - “Interestingly”. [3/4], rare
 - Inline “(i), (ii)” (0.2–1.3). [4/4]
 - **ERA-2019/20 (SC, JMP; absent in AA and BB):**
 - Heavy hedging with “seem(s) to” (SC 22, JMP 33; AA 0, BB 1).
 - The recent papers hedge with “likely”, “could” and “suggest” instead.
 - **SINGLE (do not use as rules):**
 - “In a nutshell” (JMP, BB once).
 - “Let X denote” (JMP).
 - “Result #1:” headings (JMP 2018 only; gone in later versions).
 - “Related literature.” (JMP).
 - “Taken together”, “We find”, “Notably”, em-dashes (AA only; maybe Margaryan’s hand).
 - “available upon request” (JMP).
 - “To the best of our knowledge” (AA).
 - **Anti-tidiness (round-2 judge; consistent with the paragraph and sentence spread above):**
 - no neat closing one-liners that settle a point;
 - footnotes may argue step by step (“Two problems arise. First ... Second ...”);
 - some loose ordering is fine.
-

G. Defaults (decided by me; the user said to go ahead, and each can be reversed)

- **Q1 voice:** the base is the 2024–25 voice (AA, BB), plus the STABLE rules. This follows the user’s remark that his style has evolved since the JMP. ERA-2019/20 traits (heavy “seems”, paired parentheses) are used sparingly, never as a signature.
 - **Q2 run-in headings:** sentence case, as in AA.
 - **Q3:** “percent” in running text; “%” only inside parentheses and tables.
-

I. Round-3 consolidation (3 independent judges × 2 held-out sections; checked against the corpus)

Process rule, STABLE self-bias (6/6 held-out sections, rounds 1–3):

- I write 7–21% too short. Hit the target length within $\pm 5\%$.
- Fill the gap with the kinds of material he uses, never with filler:
 - a restatement of a key fact where it matters again;
 - the institutional or data detail spelled out;
 - one more sentence of interpretation;
 - a longer footnote.

Anti-tidiness. Consensus of 3/3 judges in round 3, and also of the round-2 judge, across 5 of 6 held-out sections.

What not to do:

- Do not reorder into textbook logic (“motivate first, then define, then assume”). A paragraph may do two jobs, and a detail

may come where it becomes relevant, even if that is late (e.g. the channels after the specification, or the positioning in the literature after the identification argument).

- Do not remove all redundancy. Restate the key facts (the size of the reform, the main estimate, the two distances) where they matter again. Repeat nouns rather than using pronouns.
- Do not balance paragraphs. Next to short ones, there is usually one overloaded paragraph (the setting + data + reform paragraph in intros).

How to write footnotes:

- Footnotes are uneven. Some are one line; one or two are loaded with several loosely related papers and asides ("A shortcoming of our approach (common to other, similar papers) ...", "discussing this is beyond the scope of this paper").
- Do not curate footnotes to one topic each.

Equations:

- Gloss them in one long sentence: "where X is ...; Y is ...; ...; and ϵ is the error term" (SC, BB).
- Then one sentence on the parameter of interest.
- Do not use (a)/(b) lists for the terms.

Literature:

- Cite author by author, sentence by sentence ("X (year) finds ... Y (year) shows ...").
- Do not compress several countries or studies into one tidy parenthetical list.

Linking design to challenges in the intro:

- "Three aspects of this paper allow us to ... First ... Second ... Third ..." (SC); "Our setting offers several advantages over prior work." (AA).
- [2/4] → use one such sentence to tie the design to the challenges.

Markers checked against the corpus (counts in AA, BB, SC, JMP):

- Use:
 - "unique" [4/4], "universe" [4/4], "henceforth" [3/4], "Interestingly" [3/4], "respectively" [3/4].
 - "so-called" [2/4] and "Regarding ...," [SC] sparingly.
- Do NOT use as signature (single paper or era):
 - "Generally," and "Accordingly," (SC, JMP)
 - "Organization of the Paper." (SC, JMP)
 - "Importantly," and "It is important to reiterate" (BB)
 - "One can argue" (BB)
 - "In a nutshell" (BB, JMP)

Ceiling: judges also detect his non-native slips. We do not imitate errors, so perfect indistinguishability is not the goal. The goal is his structure and habits without his mistakes.

Tension to watch: the user earlier flagged filler and repetition in our drafts as AI-like. Montalbán restates key facts, and the user decided that Montalbán wins. Still, restate only *key* facts, and never announce what comes next without adding content.

J. Round-4 consolidation (3 judges; repeated consensus across rounds 2–4)

1. **Order.** State the key finding early (intro P2, before or right after the design), then revisit it. Walk through the material as he does: paper by paper in the literature (one paragraph per key paper, "Our paper expands on / complements X by ..."), and figure by figure in the results. Late asides are fine. Do NOT rebuild the section in textbook order.
2. **No synthesis sentences** that he did not write:
 - no topic sentences such as "These results have several policy implications";
 - no closing "Overall, ..." in every section, since the corpus rate is ≤ 0.35 per 1k words;
 - end paragraphs on a fact or a plain claim.
3. **Footnotes.** Literal and terse ("See Table 1."), or long and additive ("Besides ... Additionally ... Finally ..."). Never add explanatory clauses to them.
4. ****My own tics, measured by `tools/tic_lint.py`.** ** Run it on every draft; single-occurrence flags on texts under ~1,500 words are soft.
 - Hard limits: "therefore" $\leq 0.6/1k$ (I used 2–3/1k); "Hence" $\leq 0.7/1k$; "On the one hand ... On the other hand" ≈ 0 (I used

it in 3 of 4 texts; he never did); "Recall that" 0; "Overall," $\leq 0.35/1k$.

- Must appear: "However" (0.4–1.6/1k; I used 0), "e.g./i.e." (0.5–2.2/1k), "potential(ly)" ($\geq 0.8/1k$).

5. **ERA-2019/20 traits** ("seem(s) to" hedging; condensed "X (Y) by 50% (12%)"): all judges flag their absence in SC texts. The corpus has them in SC and JMP but not in AA or BB. For the Jornada paper the default is still the AA/BB voice (G, Q1). **Open question for the user** once the blind tests end.
 6. **Expanding to reach the target length** is where I invent. Pad only with restatements of key facts and details that are in the sources (round-4 incident).
-

H. Pre-send checklist

- Argument plan and gaps written, and the user's questions answered.
 - Length within $\pm 5\%$ of the target (section I). Lengths within section C. Paragraphs vary. No paragraph under ~ 50 words.
 - Intro: "this paper" by P2–P5, literature after the results, no double scaffolds.
 - Every estimate has its unit, relative size and significance in words. Balance and nulls are terse.
 - Every robustness claim names its appendix table.
 - Only STABLE rules applied, with the defaults in G.
 - `tools/tic_lint.py` passed (no hard-limit violations).
 - Lu Han list passed, except where Montalbán wins.
 - Numbers match facts/.
-

K. Round-5 consolidation (AA2025 held out; 3 judges; 25-09-2026)

1. **Numbers for nulls stay in the tables.** For a null effect, write "precisely estimated null effects", "small and statistically insignificant" and point to the table. Give coefficients in the text only for the headline estimate or for effects that are significant. This qualifies section E: the "unit + relative size" rule applies to the effects that matter, not to every row.
 2. **Apply section C:** run-in labels inside Data and Outcome subsections ("Parental outcomes."). Do not replace them with "The first/second group of outcomes...".
 3. **Results before motivation** (4th round with this tell). A Results subsection opens with a short technical pointer (linkage, validity), goes to the first table, and gives the "why this outcome matters" literature afterwards, in its own paragraph.
 4. $\sim$ REVOKED in round 6 (overfit; 0 corpus hits for "We do so because") $\sim$ **Action, then reason, in separate sentences:** "We drop X. ... We impose this restriction because ...". Shares and caveats get their own short sentence rather than a parenthesis.
 5. $\sim$ REVOKED in round 6 ("In sum" in 1/4 papers; judges 3/3 flagged it) $\sim$ **Closing a subsection:** a stacked restatement of 2–3 sentences ("In sum, ... These patterns suggest ... Taken together, ...") is his habit. This supersedes J.2 for the *end of a subsection*, but not for topic sentences.
 6. $\sim$ REVOKED in round 6: back to Q3 $\sim$ **Q3 revised:** "%" and "percent" may be mixed within a section, as in AA. Pending user confirmation.
 7. A loose trailing fact at the end of a paragraph is fine.
-

L. Round-6 consolidation (25-09-2026)

Rule of rules: judge feedback from one held-out section is a hypothesis. It becomes a rule only when the corpus check finds it in at least 3 of the 4 units (AA, BB, SC, JMP). Round 5 broke this, and round 6 penalised it.

1. **Sentences:** keep the reason inside the sentence, joined with ", and", ";;" or "since". Do not chop explanations into short follow-up sentences.
2. **Endings:** end a subsection on the last result or on a forward pointer ("We turn to this question in the next section."). No recap paragraph.
3. **Citations:** put them in a parenthesis at the end of the claim, or write "X (year) show that...". Never a stand-alone "See, e.g." sentence.

4. **Precision:** state what the text needs in the fact sheet's own terms. Do not add extra shares, significance levels on balance rows or unit corrections just because they are available.
 5. **Length:** if I am within about 10 percent below the target, I stay short rather than pad. Padding with interpretation has been caught in rounds 4 and 6. If I must add, I add institutional or data detail, never interpretation.
 6. **Empirical strategy order:** the design in words, the equation and its "where" gloss, the estimator, and then the identifying-variation figure.
 7. **Cross-references:** a roadmap does not need "(Section 5.1)" after each part.
-

M. Round-7 consolidation (25-09-2026)

1. **Footnotes are MY decision; the fact sheet does not make it.** Target 1.7–4.3 footnotes per 1k words, with a check at the end. Put these in footnotes: construction details of variables, "results available upon request" checks, alternative specifications, demographic or contextual side facts, legal references, and the direction of a bias when it is an aside. The main text keeps the argument and the pointer.
 2. **Length: no padding of any kind** (L.5 is superseded). If a draft is short, first check whether material that belongs in footnotes is missing from the draft; footnote words do not count toward the target anyway. Otherwise leave it short.
 3. **Numbers in prose:** the headline coefficient, its significance and one verbal comparison ("about five times smaller", "slightly lower"). Other columns are described qualitatively ("This hardly affects the estimates."). Never give ranges across columns.
 4. **Data sections:** one run-in label per source or block ("School database.", "Household parental education.").
 5. **Rhythm:** stay inside the corpus ranges that `tools/tic_lint.py` now checks (median sentence 22–29 words, 6–13% short sentences, 1.9–4.0 semicolons per 1k, 1.7–4.3 footnotes per 1k). Do not follow a single judge's advice about sentence length.
-

N. Round-8 consolidation (25-09-2026)

1. **What goes in footnotes** (in addition to M.1): legal hierarchy and law numbers; speculative explanations of a pattern ("One possible reason is..."); bare "See X (year)." pointers; grab-bag notes that mix a figure pointer with a rule.
 2. **Banned polish** (3/3 judges, 0–1 hits in the corpus): literary inversions ("Only in ... does ..."); "that is," glosses; one-line interpretive closers at the end of a paragraph.
 3. **Era traits** ("seem(s) to", "is associated with", paired "X (Y)" figures, "Overall, there are N findings"): SC and JMP only. They are NOT applied by default. Decision pending with the user (J.5).
 4. **J.5 decided (user: "lo que veas", 25-09-2026):** voice = AA/BB 2024–25. ERA-2019/20 traits are NOT used. From round 9, blind tests hold out AA and BB sections only, so that the test measures the voice we will write in.
-

O. Round-9 consolidation (25-09-2026)

1. **Footnote rate is a CHECK, not a quota.** Every footnote must carry substance: a caveat, a benchmark, a list of exclusions, a mechanism, a legal reference, a robustness result. Never write pointer-only or source-only notes ("See Table X for the full estimates.", "The data come from X.").
2. **Placement is decided by content, never by length.** Do not move material between text and footnotes to hit a word target.
3. **Two kinds of tables:**
 - Descriptive or summary-statistics tables: walk them panel by panel with 3–5 concrete values (SKILL D, [4/4]).
 - Regression tables: the headline estimate, its significance and one comparison (M.3).
4. **Where his extra length comes from** (the legitimate way to close my 10–25% gap): signposting a paragraph's job ("Note that our period ends in 2008 for two reasons: (i)..."), re-introducing an institution or variable when it reappears, and the table walk-through above. Never extra facts, interpretation or unexplained numbers.

Blind-test log (Montalbán style)

Round	Date	Held-out text	Writer inputs	Judges	Result	Score
1	2026-09-24	JMP 2018: H1 Institutional Background, H2 Results (6.1–6.5)	fact sheets + skill v1 (AA, SC, BB only)	1	detected 2/2 (93%, 95%)	—
2	2026-09-24	JMP 2018: H3 Data + Empirical Strategy, H4 Introduction	fact sheets + skill v1 + round-1 lessons + argument plan	1	detected 2/2 (90%, 90%)	6/10

Contamination note: all JMP versions (2018, PSE Jun/Nov 2019, thesis chapter 1) are now read. Future held-out sections must come from other papers. The candidates are GENDER2019 (thesis ch. 2) and the SC2020jan-only passages, and the corpus used to build the skill must then exclude that unit.

Lessons from rounds 1–2 were checked against the corpus matrix (`style/montalban/_evidence_matrix.txt`). Several turned out to be JMP-only, e.g. the paired parentheses, "Let X denote", "Result #1" headings, "Related literature." and "In a nutshell". See the evidence column in SKILL.md.

3	2026-09-24	SC2020 Introduction (H5); BB2024 Empirical Strategy (H6)	fact sheets + skill v2 (consolidated, with evidence) + argument plan	3 independent (reference: AA2025 + JMP2019nov only)	detected 6/6 (85–90%)	4, 6, 5 → mean 5.0
---	------------	--	--	---	-----------------------	--------------------

Round 3 notes:

- Length: I again wrote 17–21% short (H5 1,703 vs 2,149; H6 3,401 vs 4,095). Across rounds 1–3, 6/6 sections were 7–21% short → STABLE self-bias.
- Consensus tells, cited by 3/3 judges in both sections and matching round 2:
 - too tidy: logical reordering, one job per paragraph, balanced paragraphs;
 - I remove his redundancy and restatements;
 - my footnotes are too neat and balanced, while his are chatty, loaded and full of asides;
 - I format notation too uniformly.
- Two judges also used his non-native slips as tells. We will not imitate those, so this sets a ceiling on how indistinguishable the text can be.
- Judge-proposed markers were checked against the corpus:
 - stable: "unique" 4/4, "universe" 4/4, "henceforth" 3/4, "Interestingly" 3/4;
 - two papers: "so-called", "is the error term";
 - paper-specific, NOT rules: "Generally,", "Accordingly,", "Organization of the Paper." (SC, JMP); "Importantly,", "It is important to reiterate", "One can argue", "beyond the scope" (BB only); "In a nutshell" (BB, JMP).

Round 4, writing notes:

- First drafts were again short (H7 –9%, H8 –7%). I expanded them to within $\pm 5\%$.
- While expanding, I added two interpretations that are NOT in the facts ("these magnitudes are not negligible..."; "the inter-district reform increased distance in all groups", which is wrong, since the small municipalities had no inter-district change). I caught both before judging and replaced them with fact-based restatements.
- Lesson: expanding to reach length is exactly where invention creeps in. Pad only with restatements and details that are in the facts.

4	2026-09-24	BB2024 Introduction (H7); SC2020 Robustness §6 + Conclusion (H8)	fact sheets + skill v2 + round-3 rules + $\pm 5\%$ length	3 independent (same protocol as round 3)	detected 6/6 (72–88%)	6, 6, 6 → mean 6.0 (round 3: 5.0)
---	------------	--	---	--	-----------------------	-----------------------------------

Round 4 notes:

- Small improvement: score 5.0 → 6.0; confidence down from 85–90% to 72–88%.
- Consensus tells (3/3 judges, repeating rounds 2–3):
 - (1) I reorder into textbook sequence. He states the key finding early (P2), then walks paper by paper and figure by figure, with late asides.

- (2) I add synthesis and topic sentences (“Overall, ...”, “These results have several policy implications”).
- (3) I “improve” footnotes by adding explanatory clauses.
- (4) I cut his hedging and his condensed parentheses. Both are ERA-2019/20 traits (SC, JMP).
- The new linter `tools/tic_lint.py` (corpus rates per 1k) quantifies MY tics across H5–H8:
 - “therefore”: 15 in mine vs 2 in the originals;
 - “On the one hand/other hand”: 3 of 4 of my texts, 0 in the originals;
 - “Overall,”: 4 vs 1;
 - “Hence”: 5 vs 3;
 - “However”: 0 in all four of mine vs present in the originals.
- Single-occurrence flags on short texts are noise; the originals also trip them.

5	2026-09-25	AA2025 Data §3 (H9); AA2025 §5.3 long-term outcomes + parents (H10)	fact sheets + skill v2 + round-3/4 rules; length hit on the first draft (1,808/1,800; 986/1,000)	3 independent (reference: SC2020may + JMP2019nov; AA held out)	detected 6/6 (85–90%)	H9: 7, 7, 6; H10: 6, → mean 6.3 (round 4 6.0)
---	------------	--	---	---	-----------------------	---

Round 5 notes:

- Length is now under control: both texts were within $\pm 2\%$ on the first draft.
- Consensus tells (3/3 judges):
 - (1) I did NOT use run-in labels in 3.3, although SKILL C says to. This violates my own rule; the original had “Household school choice patterns.” etc.
 - (2) I put the motivation up front (the well-being literature in P1). He goes straight to Table 8 and gives the motivation later, in its own paragraph. This is the 4th round with this tell.
 - (3) I put too many numbers in the text. For NULLS he reports no coefficients (“precisely estimated null effects”, “small and statistically insignificant”). He gives magnitudes only for the headline or for non-null effects (e.g. the 0.03–0.06 item range).
 - (4) I merged action and justification into one sentence. He writes the action, then “We impose this restriction because...”, and puts caveats in separate short sentences (“Non-citizen children constitute 10.4% of the sample.”).
 - (5) He stacks conclusions (“In sum... These patterns suggest... Taken together...”); I compressed them. Same as rounds 3–4.
 - (6) He mixes “%” and “percent” within a section; my Q3 default normalised it. AA does mix them.
 - (7) Loose trailing sentences (“The data also show that the region of Madrid has a rich school supply.”).
- Ceiling: all 3 judges named his typos and non-native slips as decisive (“just above than”, “linage”, an unfinished sentence). We do not imitate errors (SKILL I).
- 5.2 vs 5.4: the text of the original says 5.4, and the fact sheet flagged that the table says 5.167. I used the table value. The judges read it as a slip. Not a style issue.

Protocol change before round 6 (user: “sin trampas”, 25-09-2026):

- Rounds 1–5 fact sheets leaked author choices: “(was in a footnote)” tags and the footnote count. That was a mild advantage for me.
- From round 6 the fact sheets have neither. They also do not mark which numbers the original states in the text. That information lives only in `H_manifest.md`, which the writer never reads.
- Scores from round 6 onward are therefore NOT strictly comparable with rounds 1–5, and a drop may just reflect the stricter protocol.
- Standing rules:
 - the writer never opens `_original.txt`, `H_manifest.md`, `_KEY_.json` or `judge/`;
 - no deliberate typos or slips to fool the judges;
 - no editing of a text after seeing the judges’ verdicts.
- Round-6 contamination (disclosed): the fact-sheet agent’s hand-back report mentioned three text choices of the original H11: it says “5 meters closer” for a Table 4 gap, “3 percent” for 0.032 SD, and “close to zero” for the six balance estimates. I write only from the fact sheet, but I cannot unsee these. Treat H11 as mildly contaminated.

6	2026-09-25	AA2025 §4 Empirical Strategy + Internal Validity (H11); §5 opening + §5.2 (H12)	stricter fact sheets (no footnote tags or counts) + skill v2 + K	3 independent (reference SC2020may + JMP2019nov)	detected 6/6 (H11 85%; H12 75%)	H11: 6, 6, 6; H12: 6, 6.5 → mean 6.25 (no strictly comparable: stricter protocol)
---	------------	--	--	--	------------------------------------	--

Round 6 notes: OVERFITTING of the round-5 lessons (K)

- K.5 (stacked closing "In sum... Taken together..."): 3/3 judges called it an AI wrap-up. Corpus: "In sum" appears in 1 of 4 papers, and a stacked close is not a rule. I derived K.5 from ONE held-out section, which breaks the user's rule that a style rule needs evidence across several papers. → REVOKED.
- K.4 (action, then reason in a separate sentence): it produced chopped sentences ("We do so because..."), which 3/3 judges called an imitation device. Corpus: "We do so because" occurs 0 times. → REVOKED. His real habit is long sentences chained with ", and" or ";;" and the reason inside the same sentence.
- K.6 (mix "%" and "percent"): all 4 papers mix them across the paper, but mixing inside one paragraph reads as deliberate (3/3). → Back to Q3: "percent" in running text, "%" only inside parentheses.
- K.2 and K.3 held: run-in heads were fine; the judges now flag other things.
- New consensus tells (3/3):
 - (a) Padding sentences that interpret or restate ("In our setting, the extra point does not bring...") were flagged in H12. These were exactly the sentences I added to reach the target length. This is the third time padding has been caught (J.6).
 - (b) A stand-alone "See, e.g., X; Y; Z." sentence: 0 occurrences in the corpus. Citations go in a parenthesis at the end of the claim, or as "X (year) show that...".
 - (c) I "correct" the fact sheet to be more precise ("3% of a standard deviation", 30/30/40 split, significance levels on balance rows). That is fact-sheet dumping.
 - (d) Order: equation, then estimator, then identifying-variation figure. I put the figure first.
 - (e) Mechanical cross-reference parentheticals "(Section 5.1)" in roadmaps.
- Meta-lesson: judge feedback from one held-out section is a HYPOTHESIS. It becomes a rule only after the corpus check across AA, BB, SC and JMP, as the user demanded from the start.

7	2026-09-25	SC2020may §3 Data and Summary Statistics (H13); BB2024 §4.2–4.4 main results and robustness (H14)	strict fact sheets + skill with K revised + L	3 independent (reference AA2025 + JMP2019nov)	detected 6/6 (85–92%)	H13: 5, 6, 5; H14: 5, 6 → mean 5.5 (round 6.25)
---	------------	---	---	---	-----------------------	---

Round 7 notes: the drop is on me

- DECISIVE (3/3): I wrote ZERO footnotes. SKILL C already required 1.7–4.3 per 1k words (the H13 original has 8). Once the fact sheets stopped tagging footnotes (protocol change), I stopped deciding footnotes myself. The stricter protocol exposed a rule that I never applied on my own.
- 3/3: over-reporting numbers in H14. These were exactly the numbers I added to reach the target length (L.5 allowed "data detail"). Padding has now been caught in rounds 4, 6 and 7, including padding with facts.
- 3/3: run-in labels in data sections ("School database.", "Applications for ..."). I used one; he labels each source. K.2 was only half applied.
- Mixed advice, checked against the corpus (script in the round-7 commit of tic_lint):
 - Judges said "short verdict sentences, not semicolon chains". Corpus: 6–13% of sentences under 12 words, median 22–29, semicolons 1.9–4.0/1k. My H13 (median 28, 7% short, 3.0/1k) is inside the range, so this advice is overfit on one section. Keep the corpus ranges.
 - "In conclusion": 0–0.05/1k in the corpus, so NOT a rule. "We now": ≤0.17. "Note that", "Moreover", "Hence", "As a result": 0.1–0.5/1k each, i.e. at most one each per section.
 - Judges asked for his limitations block ("We find two limitations in the data. First... The second...") and for limitations placed at the END of the Data section. Single observation; hypothesis only.
- Ceiling again: all 3 judges cite his slips and the source's factual errors (birth year, dates), which I silently corrected. We keep correcting errors.
- Round-8 leak (disclosed): the fact-sheet agent's report listed the footnote LINE numbers (4 per section). My corpus target (1.7–4.3 per 1k) already implies about 2–4, so the leak is mild.

8	2026-09-25	AA2025 §2 Institutional Background (H15); SC2020may §5.1 first results (H16)	strict fact sheets + skill with M (footnotes decided by me, no padding: H16 left 16% short)	3 independent (reference BB2024 + JMP2019nov)	detected 6/6 (H15 75–85%; H16 90–92%)	H15: 6, 6, 6; H16: 5, 5 → mean 5.5
---	------------	--	---	---	---------------------------------------	------------------------------------

Round 8 notes

- Footnotes are now present (4 and 3), and the judges no longer flag their absence. They now flag their CONTENT:
 - he puts the legal hierarchy (LOE, Leyes Orgánicas) and speculative explanations in footnotes; I put them in the text;

- bare "See X (year)." citations also go in footnotes (AA 1, BB 1 in the corpus).
- H16 (SC, 2020): 3/3 judges asked for "seem to", "is associated with", paired parentheses and "Overall, there are two findings... First... Second...". Corpus check: "seem(s) to" appears in SC 22 and JMP 41 but AA 0 and BB 0; "is associated with" only in SC (8). These are ERA-2019/20 traits. My default voice is AA/BB by design (G.Q1), so part of the H16 penalty is era mismatch and not a style failure. That is the open question J.5 for the user.
- Consensus polish tells (3/3), not era-specific: interpretive one-line closers ("This implies that...", "In our setting, school rankings should be interpreted..."); the literary inversion "Only in 2015 and 2016 does..."; "that is," glosses (corpus: 1 occurrence in 4 papers); reordering into logical blocks.
- Mixed advice rejected by the corpus check: "Our paper" vs "This paper" is paper-specific (BB and SC use "Our", AA and JMP "This"); "N findings to extract" has 0 hits; the placement of robustness notes (text vs footnote) differs between BB (footnote) and SC (text).
- Trend over 8 rounds: 4.0 (r1 est.), 6.0, 5.0, 6.0, 6.3, 6.25, 5.5, 5.5. Since round 6 the protocol is stricter. The last 4 rounds sit on a 5.5–6.3 plateau, and every judge still names his slips and factual errors as the strongest tell. Diminishing returns.

9	2026-09-25	AA2025 §5.1 household school choice (H17); BB2024 §2 background, data and descriptives (H18)	strict fact sheets + skill with M, N + voice AA/BB decided	3 independent (reference SC2020may + JMP2019nov)	detected 6/6; H17 only 72–75% (lowest so far)	H17: 7, 7, 7; H18: 5, 5 → mean 6.17
---	------------	--	--	--	---	-------------------------------------

Round 9 notes

- H17 is the best section so far (7/7/7, confidence 72–75%). The judges' tells were mostly his slips plus two things of mine: (a) a FILLER footnote ("See Table 6 for the full set of estimates..."); (b) I moved the correlation benchmark from a footnote into the text.
- H18 (5/6/5): 3/3 judges said
 - my text was compressed: I was 23% short, and "originals are 15–25% wordier";
 - I invented source-only footnotes ("The KS1 scores come from the DfE NPD.");
 - I moved substantive footnote material (the private-school caveat, the exclusion list) into the text, which I did in order to add length;
 - I dropped the panel-by-panel walk through Table 1 (summary statistics).
- OVERCORRECTION diagnosis (third time):
 - M.1's footnote rate became a QUOTA, and I manufactured filler notes to meet it.
 - M.3 ("headline number only") was written for regression tables, but I also applied it to a DESCRIPTIVE table. SKILL D (Data, [4/4]) already says to walk summary statistics in prose with 3–5 facts.
 - I moved material between text and footnotes to fix LENGTH, not because of what it said.
- Stable self-bias confirmed across rounds 3–9: I am systematically 10–25% short without padding. His extra words come from signposting ("Note that", "First and foremost", "for two reasons: (i)..."), re-introducing things, and walking tables panel by panel. They do not come from extra facts.